Artificial Intelligence and Machine Learning

Logistic Regression

Lecture 2: Outline

- Linear Regression (Review)
- Logistic Regression (Classification)
- Optimization

Recap

Design your model

- Input scalar linear model (line fitting)
- Fitting polynomials (synthetically designing features from a one-dimensional input)

$$\{(\mathbf{x}_i, y_i)\}_{i=1}^N, \mathbf{x}_i \in \mathbb{R}^n, y_i \in \mathbb{R}$$

Design your loss function

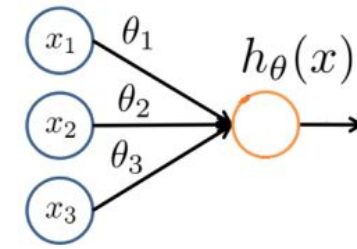- We used mean squared error loss throughout

Finding optimal parameter fitting

- Closed form solution to the linear least squares?
- Why is it linear least squares?
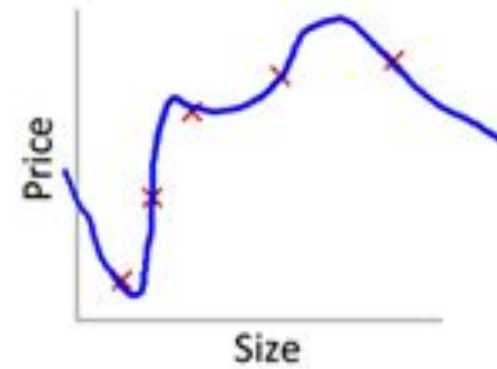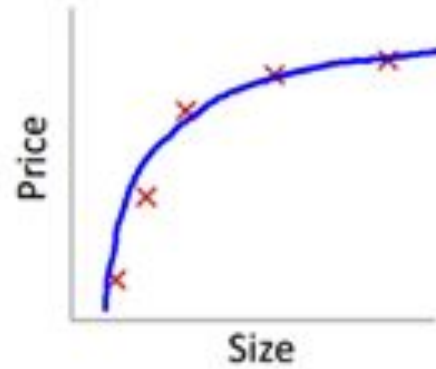- Solution is closed form

# Logistic regression or Classification

Regression VS classification

$$x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{bmatrix}$$

- Regression (linear and polynomial): for prediction

- Classification:

Regression VS classification

Income prediction -> regression

Male or Female -> classification

House price -> regression

Spam detection -> classification

Image recognition -> classification

Logistic Regression vs linear regression

| Linear Regression	Logistic Regression
For Regression	For Classification
We predict the target value for any input value	We predict the probability that the input value belongs to the specific target
Target: Real Values, continuous values	Target: Discrete values
Graph: Straight Line	Graph: S-curve

Classification

$$x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{bmatrix}$$

Binary Classification

Multi-class Classification
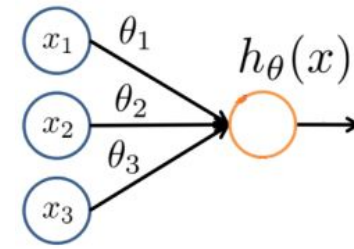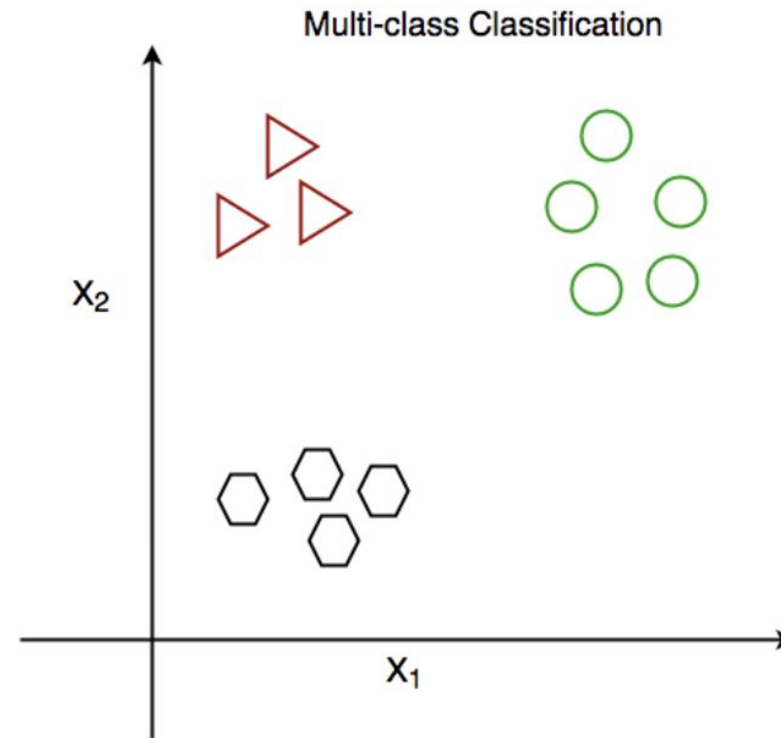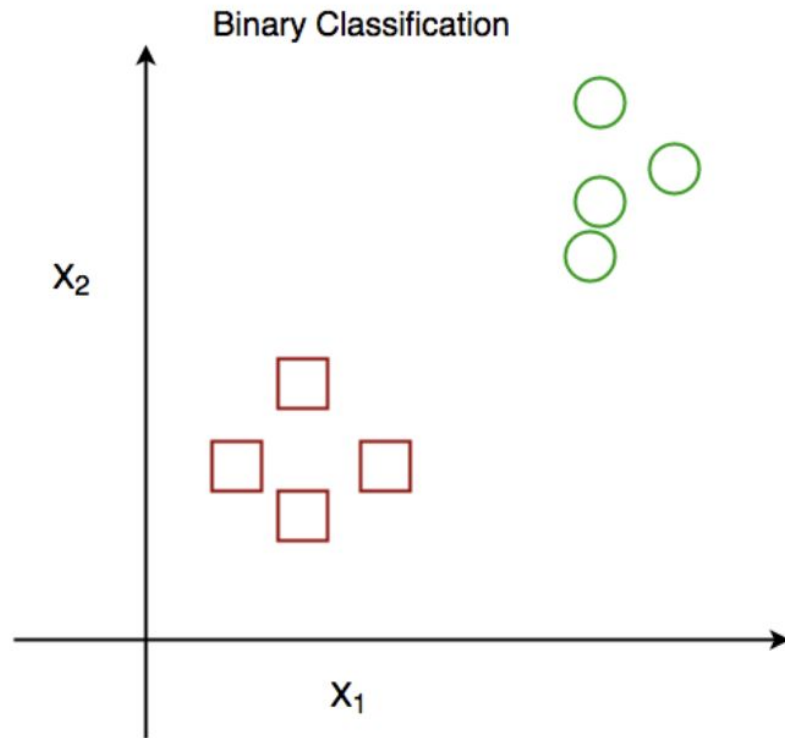
Logistic Regression

- Regular vs Fraudulent transaction
- Spam vs Non-spam emails
- Benign vs Malignant tumors
- Rising vs Falling stocks

$$\{(\mathbf{x}_i, y_i)\}_{i=1}^{N}, \mathbf{x}_i \in \mathbb{R}^n, y_i \in \{0, 1\}$$

Classification

Logistic Regression

Despite the name, logistic regression is a

classification algorithm

Logistic Regression is a linear model with a "special function" that helps us use this linear model for classification

Classification
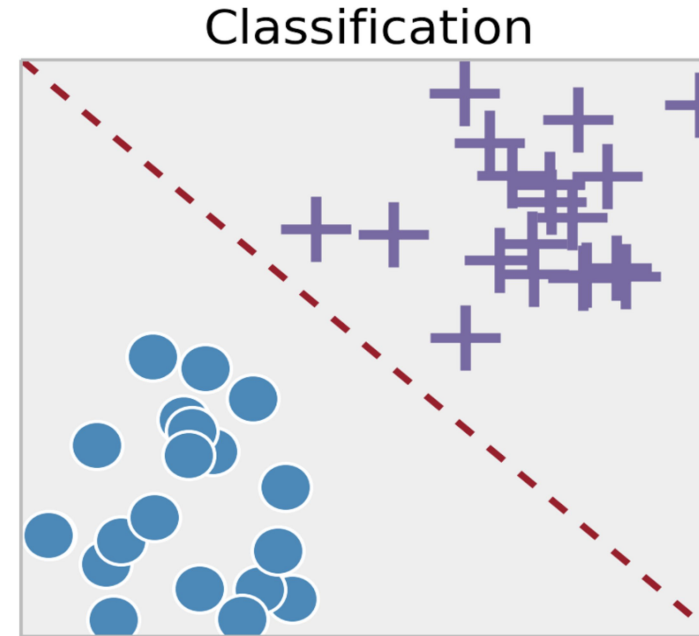
$$\{(\mathbf{x}_i, y_i)\}_{i=1}^{N}, \mathbf{x}_i \in \mathbb{R}^n, y_i \in \{0, 1\}$$

Linear Model in Disguise

$$\hat{y} = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{w} = [w_0, w_1, \cdots, w_m]^T$$

$$\mathbf{x} = [1, x^1, \cdots, x^m]^T$$

Classification
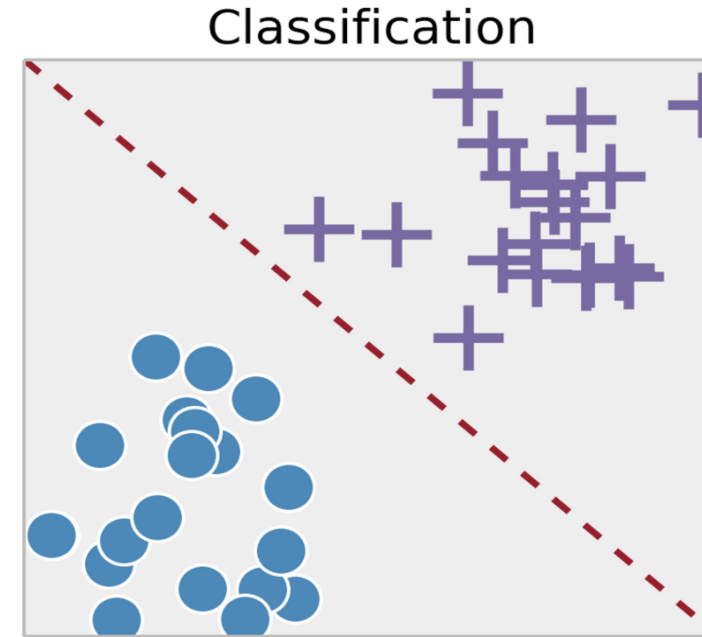
$$\{(\mathbf{x}_i, y_i)\}_{i=1}^N, \mathbf{x}_i \in \mathbb{R}^n, y_i \in \{0, 1\}$$

Challenge:

$$\hat{y} = \mathbf{w}^T \mathbf{x}$$

- The above equation predicts continuous outputs.

- Unbounded output: There is no natural constraint on the output, and prediction could be any real number

- Intuitively, it also doesn't make sense for $\hat{y}$ to take values larger than 1 or smaller than 0 when we know that y $\in$ {0, 1}.

Solution

- In logistic regression, we define the problem as follows:

  - Instead of just predicting the class, give the probability of the instance being that class

$$\hat{y} = p(y \mid \boldsymbol{x})$$

- Thus we need a function that transforms the output into a probability distribution.

$$\hat{y} = \sigma\left(\mathbf{w}^T \mathbf{x}\right)$$

Sigmoid Function

$$\sigma(z) = \frac{1}{1 + exp(-z)}$$

$$\lim_{z \to \infty} \sigma(z) = 1$$

$$\lim_{z \to -\infty} \sigma(z) = 0$$

- Assume a threshold:
  - Predict y=1 if $\sigma(z)$ >= 0.5
  - Predict y=0 if $\sigma(z)$ < 0.5
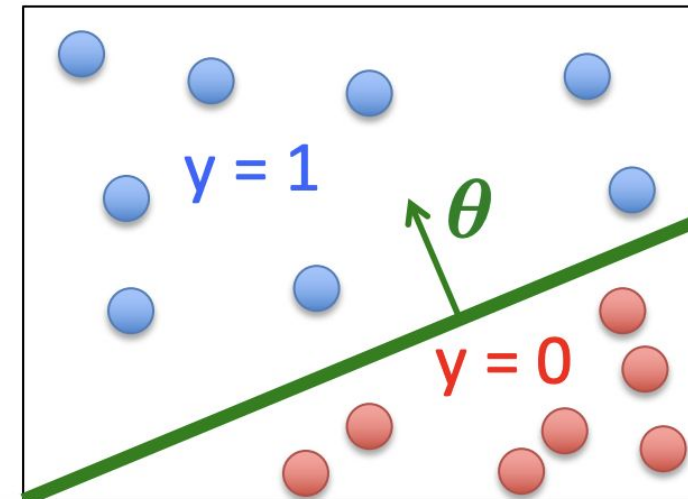
Cost Function

- We want to minimize the discrepancy between our model hypothesis and the observed label.

What type of loss to use?

MSE?
$$J = \frac{1}{N} \sum_{i=1}^{N} (y_i - \hat{y}_i)^2$$

- Not the best when it comes to classification

- Leads to suboptimal results

- Not ideal for probability output

Binary Cross Entropy Loss
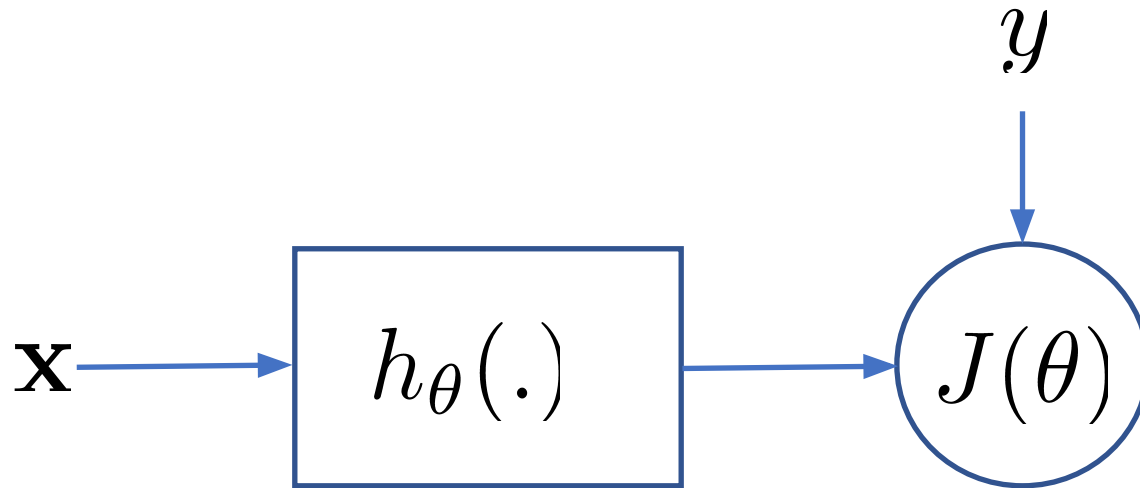
$$J(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^{N} \text{compare}(y_i, \sigma(\mathbf{w}^T \mathbf{x}_i))$$

$$J(\mathbf{w}) = -\frac{1}{N} \sum_{i=1}^{N} y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))$$

Intuition of Cost Function

- Cost of a single instance:

$$\text{cost}\left(\sigma(\mathbf{w}^T\mathbf{x}), y\right) = \begin{cases} -\log(\sigma(\mathbf{w}^T\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \sigma(\mathbf{w}^T\mathbf{x})) & \text{if } y = 0 \end{cases}$$
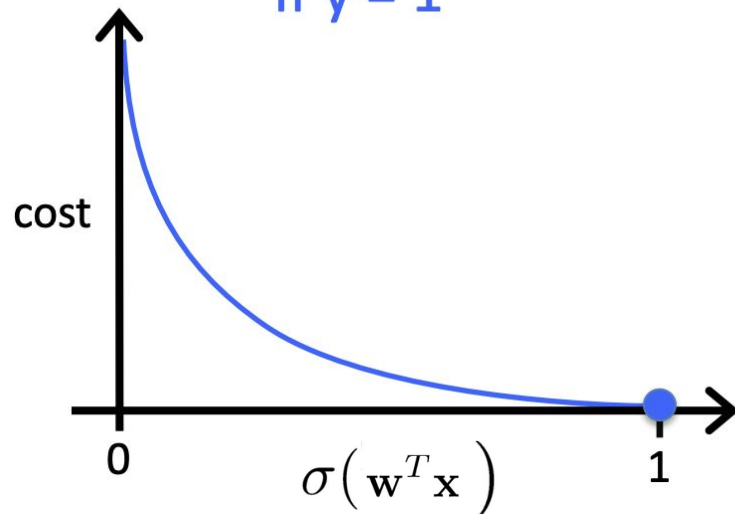
Aside: Recall the plot of log(z)
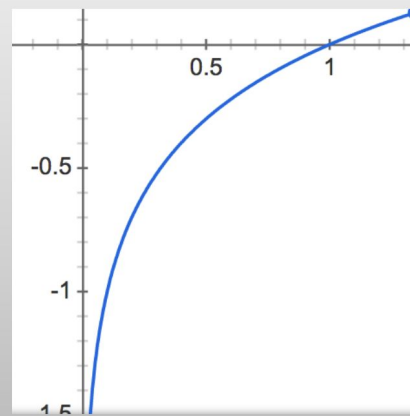
Intuition of Cost Function

- Cost of a single instance:

$$\text{cost}\left(\sigma(\mathbf{w}^T\mathbf{x}), y\right) = \begin{cases} \boxed{-\log(\sigma(\mathbf{w}^T\mathbf{x})) \quad \text{if } y = 1} \\ -\log(1 - \sigma(\mathbf{w}^T\mathbf{x})) \quad \text{if } y = 0 \end{cases}$$

If y = 1

cost

$\sigma(\mathbf{w}^T\mathbf{x})$

0   1

If y = 1

- Cost = 0 if prediction is correct
- As $\sigma(\mathbf{w}^T\mathbf{x}) \to 0, \text{cost} \to \infty$

- Captures intuition that larger mistakes should get larger penalties
  - e.g., predict $\sigma(\mathbf{w}^T\mathbf{x}) = 0$, but y = 1

Intuition of Cost Function

- Cost of a single instance:

$$\mathrm{cost}\left(\sigma(\mathbf{w}^T\mathbf{x}), y\right) = \begin{cases} -\log(\sigma(\mathbf{w}^T\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \sigma(\mathbf{w}^T\mathbf{x})) & \text{if } y = 0 \end{cases}$$
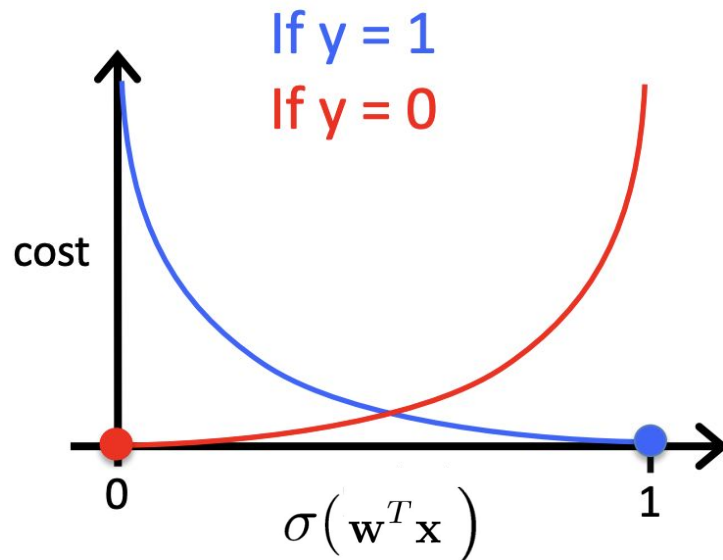
If y = 0

- Cost = 0 if prediction is correct
- As $(1 - \sigma(\mathbf{w}^T\mathbf{x})) \to 0, \mathrm{cost} \to \infty$
- Captures intuition that larger mistakes should get larger penalties

If y = 1
If y = 0

cost

0   $\sigma(\mathbf{w}^T\mathbf{x})$ 1

How to find optimal Parameters?

$$J(\mathbf{w}) = -\frac{1}{N} \sum_{i=1}^{N} y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))$$

$$\frac{\partial}{\partial \mathrm{w}_j} J(\mathrm{w}) = \frac{1}{m} \sum_{i=1}^{m} \left(\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i \right) x_j^{(i)}$$

Just like before, simply take

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$$

However, this does not have a nice closed solution.

Gradient Descent

$$J(\mathbf{w}) = -\frac{1}{N} \sum_{i=1}^{N} y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))$$

$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \nabla_{\mathbf{w}} J(\mathbf{w}^k)$$

Learning rate

- We have a linear model for prediction
- For classification, we want to output a probability
- We map the prediction to probabilities with a sigmoid function
- We have a loss function (BCE) to compare models